Fine Wine Investment Prediction: A Multi-Modal Deep Learning Approach Incorporating Subregional, Vineyard-Specific Insights, and Market Dynamics

Ignacio Estrada Cavero
Cornell University
ignacioec31@gmail.com

March 6, 2025

Abstract

Fine wine has emerged as a compelling alternative investment, known for its stable returns and low correlation with traditional markets. This proposal outlines a research plan to develop a multi-modal deep learning model for predicting fine wine investment performance. The model will integrate sub-regional and vineyard-specific features with market dynamics data to forecast an "investment score" or expected return for fine wines. We identify gaps in existing wine valuation methods—namely the lack of granular terroir-level analysis and the under use of advanced machine learning. Our approach will leverage diverse data (auction prices, weather patterns, expert reviews, images) to capture the complex drivers of fine wine prices. Key research objectives include assessing whether incorporating vineyard-level characteristics and climate trends improves prediction accuracy. We will design and evaluate a deep learning framework (combining TabNet, LSTM, CNN, and BERT modules) against traditional models and alternative assets. The proposal addresses data challenges, ethical considerations (such as market manipulation), and the potential impact on wine markets. By fusing oenology and financial data science, this research aims to contribute a novel tool for investors and enrich the academic literature on alternative asset modeling.

1 Introduction & Background

Fine wine has become an important investment asset class in recent decades, attracting investors for its combination of strong returns, tangibility, and portfolio diversification benefits. Unlike stocks or bonds, fine wines are physical commodities with limited supply, often yielding price growth independent of financial market swings. For example, an index of top investment-grade wines (Liv-ex Investables) appreciated by over 2,000% from 1988 to 2024, translating to roughly 10% annual growth. Such performance, coupled with relatively low volatility and drawdowns during crises, has positioned fine wine as a "passion asset" comparable to art or rare coins in high-net-worth portfolios. Academic and industry analyses have found that including fine wine in portfolios can improve risk-adjusted returns and hedge against stock market downturns. The growing financial significance of wine necessitates more sophisticated valuation and prediction methodologies.

1.1 Existing gaps in fine wine valuation:

Traditional approaches to wine pricing and investment value have limitations. Many studies employ hedonic pricing models, regressing wine prices on attributes like vintage, critic score, region, and producer reputation. While informative, these models often operate at a coarse regional level (e.g. Bordeaux vs. Burgundy) and may overlook sub-regional nuances such as specific village appellations or individual vineyards that can dramatically influence price. Industry benchmarks (like Liv-ex indices) aggregate broad categories of wines, which obscures the micro-level differences that collectors consider. There is a need for more granular analysis that captures how a wine's sub-appellation, vineyard characteristics, and even parcel-level terroir contribute to its investment potential. Moreover, advanced machine learning (ML) and deep learning techniques remain underutilized in this domain. Most valuation models are linear or tree-based and do not fully exploit the rich unstructured data available (such as textual tasting notes or images of bottle condition).

1.2 Need for sub-regional and vineyard-specific analysis:

Fine wine markets, especially in classical regions like Bordeaux and Burgundy, are highly stratified by classification and origin. A bottle's value can hinge on nuances like which side of a road the vineyard lies on or the reputation of a specific hillside. For instance, Burgundy's appellation system is extremely granular: the entire La Romanée Grand Cru vineyard is only 0.8 hectares, yielding about 4,000 bottles a year, and commands an average price over $7,000 per bottle. Such scarcity of top terroirs exemplifies how vineyard-specific factors drive astronomical prices. Bordeaux, by contrast, has the 1855 Classification where châteaux (producers) are ranked; First Growth estates (e.g. Château Lafite) still trade at significantly higher prices than Second Growths, reflecting the enduring impact of this historic quality ranking. These examples underscore that subregional pedigree and legal classifications (AOC – Appellation d'Origine Contrôlée – in France) are critical to fine wine valuation. However, existing pricing models rarely incorporate variables at the vineyard level (e.g. slope, soil, microclimate) or the nuanced distinctions between neighboring appellations. This research posits that integrating such fine-grained data, along with broader market indicators, will enhance the prediction of a wine's future investment performance.

In summary, as fine wine solidifies its status as a serious investment asset, there is both an opportunity and a pressing need to develop advanced, data-driven valuation methods. We aim to fill this gap by proposing a multi-modal deep learning approach that leverages vineyard-specific insights and market dynamics to predict fine wine investment outcomes more accurately than ever before.

2 Literature Review

2.1 Fine Wines vs. Traditional Investments

A growing body of literature in finance and economics has evaluated the performance of fine wine relative to stocks, bonds, and other assets. Overall, studies indicate that fine wine can deliver competitive returns with lower volatility and low correlation to mainstream markets . For example, from 2005–2020 the Liv-ex 1000 index (tracking 1,000 collectible wines) achieved an annualized return of $\tilde{8}$.4% with far smaller drawdowns during crises than equities . During the 2008 financial crisis, this wine index fell only $\tilde{1}$0% and rebounded quickly, whereas global equities plunged over

30% . Over longer horizons, fine wine's risk-adjusted performance is even more impressive: one analysis found a 15-year Sharpe Ratio of 1.53 for fine wine, vastly higher than that of gold (0.57) or the S&P 500 (0.61) in the same period . Fine wine indices have also outpaced inflation and often matched or exceeded equity markets in nominal returns . According to the Knight Frank Wealth Report, investment-grade wine prices rose 13% in 2020 alone and about 127% over the preceding decade , highlighting strong long-term growth. These findings support the idea that wine serves as an effective diversification tool . However, most index-based studies treat fine wine as a homogeneous asset class or focus on broad regions, thus offering limited insight into intra-category variability. This motivates more detailed analysis at the subregional level to see if certain wines consistently outperform others.

2.2 Determinants of fine wine pricing

Extensive research in wine economics has examined the factors that influence wine prices, often using hedonic regression models. Key determinants identified include vintage quality, critical ratings, producer reputation, scarcity, and terroir attributes . Scarcity is particularly salient in fine wine: because top wines are produced in limited quantities (sometimes just a few thousand cases annually), demand can far exceed supply, driving prices upward over time . Moreover, as bottles are consumed, the remaining supply dwindles, creating a "drinking effect" that further increases rarity and value for older vintages. Producer reputation and brand equity also heavily impact prices . Wineries with centuries of prestige (e.g. Château Margaux in Bordeaux or Domaine de la Romanée-Conti in Burgundy) command price premiums well beyond the intrinsic cost of production. Even lesser-known wines from famous appellations benefit by association with esteemed regions . Another critical factor is expert evaluations: studies have found that critic scores (such as Robert Parker's ratings or Wine Spectator reviews) correlate strongly with auction prices. A perfect 100-point score from Parker can raise a wine's price dramatically – industry observers note it might increase market price by 30–50% almost overnight . This reflects both consumer willingness to pay for perceived quality and investors treating high scores as validation of future price potential. In summary, the literature establishes that *tangible attributes* (like production volume, vineyard site, storage condition) and *intangible attributes* (brand prestige, critic acclaim) jointly determine fine wine prices. However, traditional models often struggle to capture nonlinear interactions between these factors, which is where machine learning can contribute.

2.3 AOC classifications and legal implications

Fine wine markets are heavily shaped by formal classification systems, especially in France. The Appellation d'Origine Contrôlée (AOC) framework legally defines geographic wine appellations and often implicitly ranks them by quality tier. For example, Bordeaux's 1855 Classification ranked Médoc châteaux into First through Fifth Growths based on reputation and prices at the time. Remarkably, this 170-year-old ranking still influences prices today – First Growth Bordeaux wines trade at a significant premium over lower growths, even when other quality indicators (scores, vineyard quality) are similar . This suggests an element of status value attached to the classification. The classification is enshrined in law (for instance, it's referenced in Bordeaux's AOC regulations), making it effectively a barrier to entry – no matter how good an unclassified wine is, it cannot label itself with a classified growth status. Burgundy, in contrast, has a terroir-centric classification: its hierarchy (Grand Cru, Premier Cru, Village, Regional) is tied to the specific vineyard land, not the producer . This too has legal force; only wines made from grapes in a defined Grand Cru vineyard can use that designation on the label, which confers instant elite market positioning. The financial implications are huge: Grand Cru vineyards, by law, have lower maximum yields per hectare, limiting supply to ensure quality – a factor that supports higher prices. Investors implicitly understand these legal classifications as quality signals and supply controls. Academic work (e.g., Jones & Storchmann (2001)) confirms that the 1855 classification and Burgundy's vineyard rankings help explain price differentials better than simple quality metrics . In essence, AOC laws create market segmentation that must be accounted for in any rigorous valuation model. Ignoring whether a wine is "Grand Cru" or a classified growth would omit a key driver of its investment desirability.

2.4 Vineyard-specific characteristics

Beyond official classifications, the unique attributes of individual vineyards (the concept of *terroir*) can significantly affect both wine quality and price. Terroir encompasses soil composition, slope, altitude, microclimate, and other geographical factors that impart distinct qualities to a wine. Fine-grained empirical studies are emerging that quantify terroir's value. For instance, a hedonic analysis of vineyard land in Napa and Willamette Valley found that site characteristics like slope and soil depth had measurable effects on vineyard sale prices, independent of appellation name . Similarly, in Burgundy's Côte d'Or, adjacent plots can have dramatically different market values if one is classified as Grand Cru and the other Premier Cru, reflecting centuries of perceived quality differences tied to terroir. An illustrative example: the famed Romanée-Conti vineyard (1.8 ha) produces one of the world's most expensive wines, whereas a neighboring Premier Cru, with only subtle differences in soil and exposure, sells for a fraction of the price. The literature suggests that while terroir-driven quality is real, it often intertwines with producer reputation and classification—making it challenging to isolate the effect of vineyard specifics. Nonetheless, there is evidence that micro-terroir data (e.g. heat summation in a specific vineyard, soil pH, vine age) could improve price predictions. Research on German Riesling auctions, for example, shows that including vineyard-level variables improved the fit of price models across the quality spectrum . These findings motivate our inclusion of detailed vineyard and subregional inputs in the predictive model. By accounting for the "DNA" of each wine's origin, we hypothesize we can better forecast which wines will appreciate most.

2.5 Impact of Climate Change on Wine Pricing

Climate change is an increasingly important factor in wine economics. Warmer temperatures and shifting weather patterns are already impacting vineyard yields and wine quality in traditional regions. Recent studies have noted that in the late 20th century, rising temperatures often improved ripening and led to higher wine ratings in cool regions (thus boosting prices) . However, projections for mid-21st century indicate that many classic wine regions may become less suitable for the grapes they are famous for. A study in *PNAS* modeled global wine-growing suitability and found that by 2050, areas like Bordeaux, the Rhône Valley, and Tuscany could see significant declines in climate suitability for current varietals . At the same time, cooler areas (Germany, England, parts of North America) are projected to gain suitability . Such changes have direct financial implications: if, for example, Bordeaux's top châteaux produce lower quality or smaller yields in future due to heat and drought, their wines' investment performance may suffer relative to new rising regions. The industry is already responding – we see producers buying land in new areas (e.g., English sparkling wine boom) and experimenting with heat-tolerant grape varieties . From an investment perspective, climate volatility adds uncertainty. Greater vintage-to-vintage variation or increased risk of extreme weather (hail, wildfires) can introduce volatility into wine prices. Some analyses even attribute part of the recent surge in fine wine prices to climate concerns; as global warming threatens future supply, investors may be "pricing in" expected scarcity . On the flip side, improved technology and adaptation strategies (irrigation, canopy management, etc.) could mitigate some climate risks, but often at increased cost. The literature to date underscores that any model of long-term wine investment should include climate trends or at least proxy variables (like vintage weather data). In this proposal, we plan to integrate historical weather data for each wine's region and vintage to capture the effect of growing-season conditions on quality and thus on price trajectory. This builds on the pioneering work of Ashenfelter (1995) who famously predicted Bordeaux wine prices using weather variables. By updating such approaches with modern data and climate forecasts, we can better account for market dynamics under climate change in our investment predictions.

In summary, the literature reveals: (1) Fine wine has demonstrated strong absolute and risk-adjusted returns, meriting serious investment analysis; (2) Price drivers are multifaceted, including rarity, reputation, critical acclaim, legal classification, and terroir specifics; (3) There is a research gap in leveraging detailed subregional data and multi-modal inputs (text, images) for price prediction; and (4) External factors like climate change and broader market cycles need to be incorporated into models to future-proof the predictions. Our research will build directly on these insights, using advanced deep learning to synthesize them in a predictive framework.

3 Research Questions & Objectives

Drawing from the background and literature gaps, this study is structured around several key research questions (RQs) and objectives:

3.1 Sub-regional and Vineyard-Specific Granularity

3.1.1 Research Question

Can incorporating sub-regional and vineyard-specific data significantly improve the accuracy of fine wine investment predictions compared to models using only regional or aggregate data?

3.1.2 Objective

Develop a predictive model that includes detailed geographical and terroir features (e.g., appellation, vineyard, soil/climate metrics) and test whether these fine-grained inputs increase prediction performance (lower error, higher R^2) versus baseline models without them.

3.2 Multi-modal Deep Learning Effectiveness

3.2.1 Research Question

How effective is a multi-modal deep learning approach in forecasting fine wine prices or returns, relative to traditional statistical models and simpler ML models?

3.2.2 Objective

Design a multi-modal neural network that ingests heterogeneous data (tabular attributes, time-series trends, textual notes, images) to output an investment "score" or expected ROI for each wine. Compare its performance to baseline methods like hedonic regression, Random Forest, or XGBoost on the same prediction tasks.

3.3 Modalities and Feature Importance

3.3.1 Research Question

Which modalities and features contribute the most to predicting a wine's investment value?

3.3.2 Objective

Through model interpretability techniques (e.g. attention weights, feature importance in TabNet), identify the relative importance of factors such as vineyard attributes, critic sentiment (from text), past price momentum, and market indicators. Test hypotheses such as "expert scores and vineyard reputation are the strongest predictors of price appreciation" versus "broader market dynamics dominate short-term price movements."

3.4 Adaptation to Macro-Level Market Dynamics

3.4.1 Research Question

How do market dynamics (e.g., overall fine wine index trends, economic indicators, climate events) modulate fine wine investment performance, and can our model adapt to these dynamics?

3.4.2 Objective

Incorporate macro-level time-series data (like wine index levels, equity market indices, interest rates, vintage variation due to climate) to capture broader market effects. Evaluate the model's predictive accuracy across different market regimes (bull vs. bear markets for wine) and assess if including these dynamics improves forecasts of downturn risks or boom periods.

3.5 Hypothesis

We hypothesize that:

1. **Sub-regional detail improves predictions** – wines from top-tier vineyards will be identified as having higher future returns than lesser sites, even within the same region, and the model using these details will outperform one that only knows the region name.

2. **Multi-modal integration outperforms any single data source** – a network combining text, image, and structured data will yield lower error in price prediction than models that use only one modality (e.g., just numeric features or just text sentiment).

3. **The model will capture "hidden gems"** – by analyzing diverse data, it may flag undervalued wines (e.g., from up-and-coming producers or emerging regions) that traditional models or indices might overlook, thus demonstrating commercial insight.

4. Including **market and climate variables** will make the predictions more robust year-to-year, indicating, for instance, that a string of great vintages or a global economic boom will positively shift most wines' outlook, whereas an economic crisis or poor vintage will be correctly associated with lower short-term returns.

By pursuing these questions and objectives, the research aims not only to answer specific technical queries but also to create a practical tool for investors. The ultimate goal is to produce a validated model that can predict fine wine investment performance with greater nuance and accuracy, thereby informing portfolio construction in this alternative asset class.

4 Methodology

To address the research questions, we will employ a multi-phase methodology, encompassing data collection, model development, and evaluation. The approach is inherently interdisciplinary, merging methods from finance (time-series analysis, Monte Carlo simulation), computer science (deep learning, natural language processing), and wine industry knowledge (oenological data and domain expertise).

4.1 Data Sources and Collection

We will assemble a comprehensive dataset from multiple data sources to capture the various facets of fine wine investment value:

• **Auction and Market Prices:** Historical transaction records from wine auctions (e.g., Sotheby's, Christie's, Acker) and exchange platforms (Liv-ex) will provide the dependent variable for our model – the market price or price index of fine wines over time. We plan to obtain time-series price data for a selection of investment-grade wines (e.g., Bordeaux classed growths, Burgundy grands crus, cult wines from Napa) spanning at least 10–20 years. If proprietary exchange data (like Liv-ex) is inaccessible, we will use auction price databases or published indices (e.g., Wine Spectator Auction Index). These price histories will be used both as features (for momentum/trend) and for computing target variables like future return.

• **Wine Attributes (Tabular Data):** A structured dataset of wine attributes will be compiled, including: region/appellation, producer, vineyard/cru name, grape variety, vintage year, production volume (cases produced), and critic scores (from sources like Robert Parker's Wine Advocate, Wine Spectator, etc.). Many of these features are available through wine databases (Wine-Searcher, Liv-ex's classification, or Kaggle wine datasets for reviews). The **subregional and vineyard-specific insights** are central here – for example, we will encode whether a Burgundy wine is from a Grand Cru vineyard, or the exact château and growth status for Bordeaux. We will also include any quantitative terroir metrics available (e.g., vineyard elevation, soil type classification, average vine age, if obtainable from producer tech sheets or academic GIS data on vineyards). This forms the tabular feature matrix for each wine.

• **Weather and Climate Data:** To capture vintage-specific quality drivers and climate change effects, we will link each wine and vintage to weather data from its growing season. Key variables include growing degree days, harvest time temperatures, rainfall, and extreme events (frost, drought). Data can be sourced from meteorological agencies or existing compilations (e.g., Bordeaux and Burgundy vintage weather statistics published in journal articles, or global datasets like NOAA's GSOD for weather stations near wine regions). We may also use derived indices like Parker's vintage chart ratings as a summary of vintage quality by region. These serve as numeric features indicating conditions that might influence wine quality and longevity.

• **Expert Reviews and Textual Data:** We will leverage text data from wine critics and publications. Tasting notes and reviews often contain sentiment and descriptors that are predictive of quality and thus price trajectory (e.g., a review mentioning "great aging potential" or "highly collectible" might signal future value). We will gather textual reviews for our target wines from sources like Wine Advocate, Wine Enthusiast, or community reviews (CellarTracker), where available. Natural language processing (NLP) will be applied to transform these unstructured texts into features (see Model Design below). Specifically, we'll use pre-trained language models (like BERT) to encode review text into numerical vectors representing sentiment and key themes.

• **Images:** One novel modality will be image data. Potential image inputs include photographs of the wine bottle/label and vineyard or region maps. High-end wine auctions often provide images of the actual bottle being sold; these images can reveal the bottle's condition (level of fill, label damage, coloration), which affects price for older wines. We will use any available bottle images to extract features such as fill level (ullage) detection or label recognition (to verify authenticity or rarity of formats). Additionally, if feasible, we might include satellite or topographic images of the vineyard location to provide geo-visual features (though this is exploratory). For our model, images will be processed by convolutional neural networks to derive feature embeddings that can indicate aspects like bottle condition or even visual cues of vineyard terrain.

• **Market and Economic Indicators:** To incorporate market dynamics, we will include time-series of broader indices as features. Examples are the Liv-ex 1000 index level (overall fine wine market trend), stock market indices (S&P 500, as a proxy for economic health), interest rates or inflation (as they affect alternative asset demand), and even the Knight Frank Rare Wine Index from wealth reports. These indicators, aligned by date, will allow the model to learn how external economic conditions correlate with wine price movements. For instance, we can include a feature like "1-year momentum of fine wine index" or a dummy for recession years.

All data will be merged on the appropriate keys. Each data "instance" in our dataset could be a specific wine (identified by producer-vineyard-vintage) at a specific point in time, with features from all the above categories, and the target could be a future price or return. We will need to perform extensive data cleaning (e.g., standardizing wine names, handling missing values such as missing review text or images for some wines). If certain modalities are missing for some entries (which is likely, e.g., not all wines have an image or detailed review), we will employ strategies like data augmentation or use model architectures that can handle missing modality inputs (for example, setting a mask or using an imputation network for missing images).

We will also discuss with industry partners for proprietary data access. Potential strategies include collaborating with an auction house or wine investment firm to obtain anonymized transaction data or leveraging APIs (if available) from wine databases. For instance, Liv-ex offers data services that universities can sometimes access for research. Similarly, weather data can be pulled via APIs (NOAA, Meteostat) and textual data via web scraping if not provided openly.

4.2 Multi-Modal Deep Learning Model Architecture

The core of this research is the development of a **multi-modal deep learning model** that ingests the different data types and outputs a prediction of the wine's investment performance. We define the prediction target as an **investment score or return metric** – for example, the model could predict the *annualized ROI for the next 1-3 years* for a given wine, or the probability that the wine will outperform a benchmark index. For the purposes of model training, a convenient target is the next-period price (or log-price) of the wine, which enables supervised learning via regression. We can transform that into ROI if needed (ROI = (predicted future price − current price)/current price). The model will be trained to minimize error between its predictions and the actual outcomes observed in historical data.

Model architecture: We propose a neural network architecture with separate subnetworks for each modality, which are then fused. Specifically:

• **Tabular data branch:** For structured numerical/categorical features (region, scores, production volume, etc.), we will use **TabNet**, an interpretable deep learning architecture for tabular data . TabNet employs sequential attention to select the most relevant features at each decision step, and has shown strong performance on table-style datasets. By using TabNet, we not only aim for high accuracy but also the ability to interpret which features (e.g., vineyard ranking, critic score, age) the model is focusing on . The output of the TabNet branch will be a feature embedding (a vector) summarizing the key info from the tabular inputs.

• **Time-series branch:** To capture trends and temporal dependencies, an **LSTM (Long Short-Term Memory)** network will be used on sequential data. For each wine, we can feed in a sequence such as the last T periods of price changes or index values. The LSTM, a recurrent neural network suited for sequence data, will learn patterns like momentum or mean-reversion in wine prices. It can also ingest sequences of external market indicators over time. The LSTM's hidden state after processing the sequence will form a time-series features embedding. (If certain wines have sparse transaction history, we may instead use a panel data approach or fill missing periods with index movements.)

• **Text branch:** We will fine-tune a **BERT**-based model (Bidirectional Encoder Representations from Transformers) for the wine review texts. BERT is a state-of-the-art language model that produces context-rich text embeddings . Each wine's collection of textual reviews (possibly concatenated or the most relevant one) will be input to a BERT encoder, producing an embedding vector capturing sentiment and important descriptors. This vector might capture, for instance, that a wine is described with "high acidity, needs time" (implying long aging potential), which could correlate with future price appreciation. We will likely use a pre-trained BERT (perhaps further pre-trained on wine domain text if available) and then fine-tune it on a secondary task like classifying whether a review implies "investment-grade" or not, to specialize the embeddings for our task. The final text embedding for each wine goes into the fusion layer.

• **Image branch:** For images of bottles or vineyards, a **Convolutional Neural Network (CNN)** (such as ResNet) will be employed. We may use transfer learning with a model pre-trained on ImageNet, then fine-tune on our wine images for tasks like classifying fill levels (e.g., normal, mid-shoulder, low-shoulder fills for old bottles) or detecting label flaws. The CNN's penultimate layer (after convolution and pooling) will provide an image feature vector. This could capture visual cues—e.g., a faded label and low fill indicating an old bottle in not-great condition, which might negatively impact price at auction. If we incorporate vineyard satellite images, the CNN might learn features like vegetation density or terrain type, although this is experimental. For simplicity, the image input may be primarily bottle shots from auctions.

• **Fusion and output:** The embeddings from the TabNet branch, LSTM branch, BERT/text branch, and CNN/image branch will be concatenated into a single combined feature vector. This fused representation now encompasses information from all modalities for a given wine. This vector is then passed through fully connected (dense) layers to produce the final output. The final output layer could be a single neuron (for regression of price or ROI) or a small vector (if predicting multiple targets like one-year and five-year ROI, or a classification of outperform/underperform). We will use an appropriate activation (likely linear for regression). The model's training objective will be to minimize a regression loss (such as Mean Squared Error between predicted and actual log-price or return). Here's a refined, clear, and formal phrasing of your statement in markdown format:

For optimization using Root Mean Square Error (RMSE), the loss function can be defined as follows:

$$\text{RMSE} = \sqrt{\frac{1}{N}\sum_{i,t}(y_{i,t} - \hat{y}_{i,t})^2}$$

where $y_{i,t}$ is the actual outcome for wine i at time t, $\hat{y}_{i,t}$ is the predicted outcome, and N is the total number of observations across all wines and times.

Mathematically, our multi-modal model can be described as:

$$\hat{y} = f\left(F_{\text{tab}}(x_{\text{tab}}),\, F_{\text{time}}(x_{\text{time}}),\, F_{\text{text}}(x_{\text{text}}),\, F_{\text{img}}(x_{\text{img}})\right)$$

where $F_{\text{tab}}, F_{\text{time}}, F_{\text{text}}, F_{\text{img}}$ represent feature extraction functions for tabular, time-series, textual, and image data respectively. These extracted features are combined by the model $f(\cdot)$ to produce the final prediction $\hat{y}$. We will implement training using stochastic gradient descent and backpropagation, optimizing the model parameters on our collected dataset. For cases involving pre-trained models, certain layers may remain fixed, while parameters in remaining layers are fine-tuned.

4.2.1 Training procedure

We will split the data into training, validation, and test sets, using time-aware splitting to ensure that the model is always predicting forward in time (to mimic real investment prediction and avoid lookahead bias). For example, we might train on data up to 2018 and test on 2019–2020 data. We'll employ techniques to address class imbalance if applicable (e.g., maybe only a few wines skyrocket in value, which could be seen as "positive class" in a classification sense). We'll also perform hyperparameter tuning (learning rate, network sizes, dropout rates, etc.) perhaps via Bayesian optimization or grid search, using the validation set.

Given the multi-modal nature, one challenge is ensuring each sub-network gets sufficient training. We might use a staged training: first train the TabNet and LSTM branches on their tasks (perhaps predicting prices with only those features) to give them a good initialization, and fine-tune BERT on a language understanding task for wine text (like sentiment regression). Then we'll combine and train end-to-end on the final prediction task. This approach can speed up convergence and improve performance.

We will also experiment with ablation studies: turning off one modality at a time to see the impact on performance (e.g., train model without text input to quantify how much text was adding). This directly addresses RQ3 about modality importance.

To implement this, frameworks like PyTorch or TensorFlow will be used, given their flexibility in building custom multi-input models. We may also leverage libraries for TabNet (e.g., PyTorch TabNet implementation) and transformers (Hugging Face's Transformers for BERT). The model training will be computationally intensive, so access to a GPU cluster will be important. If computation proves to be a bottleneck, we will consider simplifying the model or using dimensionality reduction on certain inputs (for instance, using a smaller language model than BERT if needed).

4.3 Mathematical Formalization of the Predictive Model

We formalize the investment prediction task as a supervised learning problem. Each wine instance is indexed by i and its age (or evaluation time) by t. We define multiple modalities of data for each wine instance as follows:

- $\mathbf{x}_i^{\text{tab}}$: Tabular features for wine i, including regional classification, grape variety, and other structured data.
- $\mathbf{x}_{i,t}^{\text{time}}$: Time-series features for wine i up to time t, such as historical prices or market index values.
- $\mathbf{x}_i^{\text{text}}$: Textual data, specifically review corpora associated with wine i.
- $\mathbf{x}_i^{\text{img}}$: Image-based data associated with wine i, such as bottle labels or related visual information.

Given these multimodal inputs, we formalize investment prediction as a supervised learning problem. Each wine instance is indexed by i, observed at discrete time intervals denoted by t. We aim to predict an outcome $y_{i,t+\Delta}$ at a future time $t+\Delta$.

Our predictive model is thus represented by the function:

$$\hat{y}_{i,t+\Delta} = f_\theta\left(\mathbf{x}_i^{\text{tab}}, \mathbf{x}_{i,t}^{\text{time}}, \mathbf{x}_i^{\text{text}}, \mathbf{x}_i^{\text{img}}\right),$$

where θ denotes the parameters of our model $f(\cdot)$.

The primary learning objective is to minimize a chosen loss function measuring the difference between the actual and predicted outcomes. For regression objectives, we define our target variable $y_{i,t+\Delta}$ as either the future wine price $P_{i,t+\Delta}$ or the log-return from time t to $t+\Delta$:

$$y_{i,t+\Delta} = \log\left(\frac{P_{i,t+\Delta}}{P_{i,t}}\right),$$

though it may also directly represent future prices $P_{i,t+\Delta}$. Using these definitions, we optimize parameters θ by minimizing the mean squared error (MSE) loss function:

$$L(\theta) = \frac{1}{N}\sum_{i,t}\left(y_{i,t+\Delta} - \hat{y}_{i,t+\Delta}\right)^2.$$

Alternatively, certain predictive objectives may be naturally formulated as classification tasks—for instance, predicting whether a wine's ROI exceeds a predetermined threshold. In such scenarios, the predicted outcome $\hat{y}_{i,t+\Delta}$ represents the probability of class membership, and the model is trained by minimizing the cross-entropy loss:

$$L(\theta) = -\frac{1}{N}\sum_{i,t}\left[y_{i,t+\Delta}\log(\hat{y}_{i,t+\Delta}) + (1-y_{i,t+\Delta})\log(1-\hat{y}_{i,t+\Delta})\right].$$

4.3.1 Interpretability and Formal Analysis

We will also formalize methods to interpret the predictions made by our model. Specifically, for the TabNet architecture, interpretability arises naturally from feature selection masks computed at each decision step. Aggregating these masks yields global feature importance scores. Formally, let a_j represent the attention weight assigned to feature j at a given decision step. The global importance of feature j can then be quantified by summing or averaging a_j across all decision steps, offering insights into which variables significantly influence predictions.

Interpreting text features similarly involves analyzing token-level importances. For instance, if a token (e.g., "premature" or "oxidation") significantly decreases predicted investment value, this interpretation aligns with known qualitative factors in wine valuation—such as indications of premature aging negatively impacting desirability.

In sum, our methodological framework emphasizes interpretability through explicit model transparency, integrating rigorous feature engineering with advanced deep learning models. The result is an interpretable predictive model, enabling stakeholders not only to forecast wine investment outcomes but also to understand the underlying reasons behind the model's predictions.

5 Comparative Benchmarking

A crucial part of our methodology is to benchmark the fine wine investment performance and our model's predictions against other assets and simpler models. This serves two purposes: (1) to validate that our model's outputs make sense in a financial context (e.g., predicted returns are reasonable compared to historical asset returns), and (2) to demonstrate the value-added of our approach over existing methods.

5.1 ROI comparisons with other asset classes

We will situate fine wine's returns in context by comparing historical return on investment (ROI) and risk metrics against real estate, equities, and art. Using data from financial markets and art price indexes (e.g., Sotheby's Mei Moses index for art), we will compute metrics like compound annual growth rate (CAGR) and volatility over matching periods. Prior research and market reports indicate fine wine's performance has been competitive; for instance, over 10-year horizons, fine wine (about 127% total gain) has outperformed many equity markets and matched or exceeded collectible art . We will update these comparisons to the present: for example, if the S&P 500 returned, say, ˜150% in the last decade (hypothetical), fine wine's 127% is in the same ballpark but with lower volatility. We will also compare to real estate (perhaps using a prime property index) and other collectibles like rare whisky (which saw huge gains in recent years). The goal is to underscore fine wine's risk-adjusted return profile. We will calculate the Sharpe ratio and Sortino ratio for fine wine vs. these assets over the last ˜20 years to see how they stack up (as touched on in the literature review) . These comparisons will be presented likely in a table or chart for clarity.

5.2 Baseline modeling with traditional methods

To benchmark our predictive model, we will implement several alternative models:

5.2.1 Linear Regression

A linear regression or hedonic model using key features (region, score, age, etc.) as a baseline for price prediction. This represents the traditional econometric approach. Its accuracy (R^2, RMSE) on the test set will be a baseline; we expect our model to exceed this if our hypotheses hold true.

5.2.2 Tree-Based Model

Tree-based ensemble models like Random Forests and XGBoost (Extreme Gradient Boosting) will be used as strong baseline ML models for tabular data. XGBoost in particular is known for its high performance on structured data and won't directly incorporate text or images (we could feed it some engineered features from those, such as sentiment score from text). We will train an XGBoost model to predict the same targets using all numeric features and some encoded categorical ones. Similarly, a Random Forest can be tried. These models have the advantages of easier interpretability and lower risk of overfitting on small data. Their results will tell us how much improvement the deep learning model offers. If the improvement is marginal, we may need to justify the added complexity; if it is large, it supports our approach.

5.2.3 Time-Series Model

We might also benchmark against classical time-series forecasting for prices, such as ARIMA or vector autoregression (VAR) that includes economic indicators. This is to see if our multi-modal RNN approach is better at capturing temporal patterns than these well-established methods.

5.2.4 Monte Carlo Simulations for Risk-Adjusted Returns

Once the model is trained, we will also use Monte Carlo simulation to evaluate investment outcomes. The idea is to simulate many scenarios of wine price movements and see how a portfolio guided by our model would perform. Specifically, we can do the following experiment: use our model's predictions to create a portfolio of wines (e.g., pick the top 10 wines that the model predicts to have highest return over next year). Then simulate random shocks or noise around the predicted returns (based on historical volatility), generating, say, 10,000 possible outcome scenarios. From these, we can compute a distribution of portfolio returns and derive risk metrics: the mean return, standard deviation, Value-at-Risk (VaR) at 5% (the 95th percentile worst loss), etc. We will compare this to a baseline portfolio (maybe an equal-weight basket of wines or the market index) under similar simulations. The Monte Carlo approach helps incorporate the uncertainty in predictions and real market variability to see if our strategy consistently adds value. For example, if our model-driven portfolio shows a higher Sharpe ratio and a lower probability of large loss across simulations, that indicates a successful prediction system. We will also test how sensitive results are to model error; if slight changes in predictions drastically alter outcomes, that's a caution on model confidence.

5.2.5 Scenario Analysis

We can test the model's predictions against known historical events. For instance, take the conditions just before the 2008 crisis or the 2020 COVID shock and see if the model would have signaled lower returns for wine (which actually happened temporarily in those periods). This isn't a traditional benchmark, but it's a sanity check aligning model output with real events.

Finally, we will evaluate simpler single-modal deep learning baselines to ensure each modality is contributing. For example, train just an LSTM on price history alone to predict future price (like a pure time-series DL model), and just a BERT model on text sentiment to predict a wine's score or price. These will likely be less accurate than the combined model, but quantifying that gap is part of our analysis for RQ3.

Through these comparative analyses, we expect to demonstrate that fine wine, as an asset, holds its own against other investments (justifying its study), and that our multi-modal model provides a measurable improvement in predicting its performance over existing benchmarks. For instance, if a Random Forest achieves an R^2 of 0.50 in explaining next-year price changes, we aim for our model to push closer to, say, 0.65–0.70 R^2. Likewise, in a trading simulation, the AI-guided selection might yield an annualized return a few percentage points higher than a naive index approach, with equal or lower volatility.

6 Expected Challenges & Solutions

Conducting this research entails several challenges, which we have identified along with proposed solutions and mitigation strategies:

6.1 Data availability and quality

Fine wine market data can be proprietary and sparse. Auction results and exchange prices are not always publicly available or might be costly. There's also the issue of missing data: not every wine will have a complete set of information (e.g., some have no Parker score or missing weather data for a given vintage).

6.1.1 Solution

We will leverage multiple sources and focus on well-documented wines. We plan to start with data that is accessible (e.g., the Liv-ex indices and publicly reported auction highs for famous wines) and augment with third-party datasets (like the Wine Spectator Auction Index or Kaggle wine reviews for text). For missing price data, we can interpolate or use related wines' indices as a proxy. We will also consider data augmentation techniques: for example, if we lack many images of bottles in poor condition, we can simulate or reuse images from similar wines. To handle missing modalities in the model, we can include mask indicators or train the model to operate with partial inputs (e.g., if no image, use an "average" image embedding). In worst-case scenarios, we might narrow the scope to wines where data is rich (like Bordeaux top growths often have plenty of data). Forming a partnership or obtaining a research data grant from a platform like Liv-ex or Wine Owners could dramatically help; we will actively pursue such collaboration for data sharing.

6.2 Heterogeneity of data and integration

Combining numerical, textual, and image data is complex. Each modality has different scales and structures, and the model must learn from all simultaneously. There's a risk that one modality (say text) could dominate or that the model has difficulty aligning information from modalities.

6.2.1 Solution

Careful architecture design (as outlined) and training procedure will mitigate this. We will normalize and preprocess each input appropriately (e.g., scaling numeric features, tokenizing text, resizing images). Early fusion vs. late fusion is a consideration – we propose late fusion (each modality processed separately then combined) which generally works well when modalities are very different. We will also monitor training for modal imbalance; for example, we can weight the losses or use multitask learning (have the text branch also try to predict critic score as an auxiliary task, for instance) to ensure each part learns meaningful representations. Another solution is to use attention mechanisms at fusion to let the model weight which modality is most informative for a given prediction. This way, if images aren't useful for certain wines, the model can effectively ignore them. Technically, the use of proven components like BERT and TabNet helps because they each handle their domain effectively; our job is mostly to get them to talk to each other in the final layers.

6.3 Computational complexity

Training a deep learning model with four sub-networks and diverse data is computationally intensive. With limited data (which is likely in this niche domain), there's also a risk of overfitting such a large model.

6.3.1 Solution

We will adopt several strategies: (a) Use transfer learning to leverage pre-trained models (BERT, ResNet, etc.) so that we don't train everything from scratch – this reduces required data and training time. (b) Implement regularization techniques such as dropout, weight decay, and perhaps early stopping based on validation loss to prevent overfit. (c) If data proves too limited, consider simplifying the model (e.g., use a smaller language model or fewer layers). We will also utilize cross-validation given the small dataset, to maximize use of data for training while still getting performance estimates. On the hardware side, we will make use of GPU acceleration and, if needed, cloud computing resources. Running the multi-modal network in parallel on multiple GPUs (one per modality, for example) might be possible to speed up training. If training time is still an issue, another approach is to train modality-specific models separately and only combine them in a second stage (which reduces the complexity of joint training).

6.4 Model evaluation and interpretability

Deep learning models can be black boxes. Convincing the finance audience of its validity requires interpretation of results (why did the model predict a certain wine will rise in value?).

6.4.1 Solution

We will incorporate interpretability from the start. Using TabNet is one conscious choice for its interpretability; we can extract feature importance for tabular data easily . For the text part, we can use techniques like LIME or SHAP to highlight words in a review that influenced the prediction. For the image part, visualization of convolutional filters or saliency maps can show if, say, the model is looking at the fill level in the bottle photo. Presenting a few case studies in the results (e.g., the model predicts Wine X will jump in price mainly because "1990 was a hot vintage with high critic scores and it's from a scarce vineyard") will help clarify the model's decision process. This is more of a presentation challenge than a roadblock, but it's one we anticipate and plan for by gathering the tools to probe the model.

6.5 Dynamic market changes

The wine market, like any market, can change due to unforeseen factors (e.g., sudden demand surge from China in 2009–2011, or tariffs on wine, etc.). A model trained on past data might not immediately account for such regime shifts.

6.5.1 Solution

While we cannot predict unknown future shocks, we can design the model to be adaptive. This might involve periodically retraining with new data (online learning) or including regime indicator features. We'll also test the model's robustness by training on one period and testing on another very different period (e.g., train pre-2010, test on the China-led boom after 2010) to see how well it extrapolates. If needed, we might incorporate external forecasts (like economic growth forecasts) to help anticipate market demand changes. In any case, acknowledging this limitation is important – the model provides a baseline prediction, but investors should still overlay current market context.

By anticipating these challenges, we increase the likelihood of a successful project. Each potential obstacle (data, integration, computation, interpretability, market shifts) has a corresponding plan to address it. The interdisciplinary nature of our team (combining data scientists and wine economists) will also help in devising creative solutions—for instance, using wine domain knowledge to fill data gaps or constrain the model (we might, for example, enforce that age has a non-linear but monotonic effect up to a point, based on wine aging science, by engineering that feature). Flexibility is key: if one approach fails, we will iterate (perhaps the multi-modal model is too ambitious for available data; we could then simplify to a bi-modal model focusing on the most impactful modalities).

7 Evaluation Metrics

To rigorously evaluate the performance of our predictive model and the outcomes of our investment strategies, we will employ a range of **metrics** addressing both the machine learning accuracy and the financial utility:

• **Root Mean Squared Error (RMSE):** This will be a primary metric for regression accuracy. RMSE measures the standard deviation of prediction errors. We will compute RMSE between predicted and actual wine prices (or returns) on the test set. A lower RMSE indicates the model predicts prices closer to the true values. We choose RMSE (over, say, MAE) because it penalizes larger errors more, which is important in an investment context (a few big miss-predictions could hurt a portfolio significantly). We will likely report RMSE in both absolute terms (e.g., in dollars or index points) and as a percentage of the mean price for interpretability.

• **R-squared (R^2):** To gauge the proportion of variance explained by the model, we will use R^2. This gives a sense of how much of the fluctuation in wine returns our model accounts for. An R^2 of 0.6, for instance, would indicate the model explains 60% of the variability in outcomes, which would be quite strong for a financial model. We will compare R^2 of our model to that of baseline models to quantify improvements.

• **Mean Absolute Error (MAE):** Though not explicitly asked, we may include MAE for completeness, as it's easier to interpret in the original units (e.g., "on average, the prediction was off by $X").

• **Classification metrics (if applicable):** If we also formulate part of the problem as a classification (for example, predicting whether a wine's annual return will be above the market median or not), we will use metrics like **ROC AUC (Receiver Operating Characteristic Area Under Curve)**. ROC AUC evaluates the true positive rate vs. false positive rate across thresholds, essentially measuring the model's ability to rank high-return wines above low-return wines. A high AUC (close to 1.0) would mean the model is good at identifying the winners vs losers in terms of investment performance. This is useful if our model output can be interpreted in a binary success context. Precision, recall, and F1-score could also be examined if we set a specific threshold (e.g., top 20% returns as "positive class"), but given the proposal's mention of ROC AUC, we'll likely focus on that for overall discrimination ability.

• **Financial performance metrics:** Beyond ML metrics, we will evaluate the **financial outcomes** of using the model:

• **Annualized ROI of model-picked portfolio:** We will simulate investing based on the model's recommendations and compute the realized annualized return. For instance, using back-testing, if the model in each year picks the top 5 wines to invest in, we track what the return of that basket is over the next year (using actual market prices) and annualize it. We'll compare this to benchmarks (market index or random selection). This measures whether the model's predictions can translate into superior returns.

• **Interpretation via Sharpe Ratio:** We assess the investment prediction model's performance using the Sharpe ratio, a measure of risk-adjusted returns. Formally, the Sharpe ratio is defined as:

$$\text{Sharpe Ratio} = \frac{E[R - R_f]}{\sigma},$$

where $E[R - R_f]$ is the mean excess return of the portfolio over a risk-free rate R_f, and σ is the standard deviation of the portfolio returns. A higher Sharpe ratio indicates superior compensation for risk, suggesting that the model effectively identifies investments with favorable returns relative to volatility. If our model accurately predicts less volatile yet profitable wines, the Sharpe ratio should notably improve compared to benchmark indices. For instance, a Sharpe ratio of 1.0 for the model-generated portfolio versus a Sharpe ratio of 0.5 for a baseline wine index portfolio would demonstrate substantial predictive value from the model.

• **Max Drawdown and VaR:** We might compute the maximum drawdown of the model's portfolio (largest peak-to-trough loss) to see if the model avoids big losses. Also, Value-at-Risk (5% VaR) can be computed from the distribution of returns either analytically or via the Monte Carlo simulation described. This gives a sense of downside risk.

• **Calibration and error analysis:** We will also assess how calibrated the model's predictions are. For instance, if the model says a wine will have +10% next year, how often is it roughly right? We can bucket predictions and see the average actual outcome for each bucket (this is a calibration curve). A well-calibrated model's predictions would align with actual averages.

• **Benchmark comparison metrics:** To directly compare models, statistical tests can be used. We can use Diebold-Mariano tests for predictive accuracy to see if the difference in error between our model and a baseline is significant. We can also report relative improvement (e.g., "our model's RMSE is 15% lower than XGBoost's RMSE").

All metrics will be computed on the held-out test set (data not seen during training) to ensure an unbiased evaluation. If cross-validation is used, we will report average metrics with standard deviations. The combination of both **statistical metrics** (RMSE, R^2, AUC) and **practical investment metrics** (ROI, Sharpe) is important. The former proves the model's technical accuracy; the latter demonstrates real-world value. It's possible a model with slightly higher RMSE could still be more useful if, say, it gets the direction of price movements right more often (hence good ROI). So we will interpret results in a balanced way.

We also plan an **error analysis** where we examine cases of large errors to see if there is a pattern (e.g., the model struggles with certain regions or with sudden market regime changes). This can highlight limitations and areas for future improvement.

In summary, success will be measured not just by how low the prediction error is, but by whether the model's usage could realistically improve an investor's decisions. If we can show, for example, that using the model one could have achieved a higher Sharpe ratio portfolio of wines than just buying the wine index or randomly picking wines, that would be compelling evidence of the model's utility.

8 Ethical & Commercial Considerations

While this research focuses on technical advancement and financial outcomes, it is important to address **ethical and commercial implications** of deploying AI in fine wine investment:

8.1 Market manipulation and fairness

An AI-driven model that identifies undervalued wines could potentially be used to exploit the market. If the model's recommendations were made public or used by a large fund, it might lead to coordinated buying of certain wines, thereby driving up their prices artificially. This could disadvantage smaller collectors and distort the market's natural price discovery. There is a risk of a self-fulfilling prophecy: e.g., the model says wine X will rise, many follow that advice and buy wine X, which indeed makes its price jump – not purely because of intrinsic value but because of the model.

8.1.1 Mitigation

We will stress that the model is a decision-support tool, not an oracle, and it should be used responsibly. Any commercial application should consider throttling the impact on the market (perhaps by broadening recommendations, not focusing all on one asset). In publishing research results, we will avoid promoting it as a "get-rich-quick" tool or giving explicit investment advice that could manipulate the market. Additionally, transparency can help: if investors know many others use a similar model, they can account for that in their strategy, reducing herding effects.

8.2 Impact on Small Producers and Market Inclusivity

If successful, the model may steer investment towards certain regions or producers that have the data and track record to be favorable. Top-tier producers with lots of historical data might be consistently picked as best investments (since they have been in the past), which could reinforce the concentration of money in famous names (like First Growths or DRC). This might mean lesser-known quality producers get overlooked, which could widen the gap between the "blue-chip" wines and the rest. However, one could argue the opposite: the model might uncover hidden gems that were undervalued, thus bringing attention and capital to smaller producers. There's an ethical dimension in ensuring the technology benefits the broader wine community, not just those already at the top.

8.2.1 Consideration

We will explore the model's picks to see if it has any bias towards established regions vs emerging ones. If a bias is found, one could adjust the model or inform users to include diversity or sustainability criteria in decisions. On a commercial level, if a wine investment fund uses this model, they might decide to also support lesser-known wineries (perhaps for ESG reasons or future growth potential) rather than purely chasing short-term gains. We will mention these possibilities in our discussion.

8.3 Transparency and Trust

Investors and wine professionals might be skeptical of a "black box" AI making recommendations on something as nuanced as wine. Ethically, providing explainable results is important so that users trust the model for the right reasons. We aim to present not just "Wine A will go up 10%" but *why* the model believes so (e.g., "it has a 100 point score, coming from a tiny Grand Cru vineyard and past similar wines have done well"). This builds trust and ensures the human decision-maker remains informed, not blindly following AI. In a commercial setting, a company using this model should maintain a human expert in the loop to validate predictions.

8.4 Data privacy and Intellectual Property

Although our data deals with wines (not personal data), some data sources might be proprietary or sensitive (e.g., a trading platform's transaction records). We must ensure any non-public data is used in compliance with terms and possibly anonymized. If we partner with a platform for data, we may have to keep certain results confidential or aggregate them to not expose sensitive information. Also, if the model incorporates text from copyrighted reviews (e.g., Robert Parker's notes), we need to respect IP – perhaps only using publicly available excerpts or obtaining permission for research use. For commercial use, licensing of such data would be required.

8.5 Commercialization and Conflicts of Interest

Should this model be commercialized (e.g., as a service for investors or part of a wine fund's strategy), there are ethical considerations around conflicts of interest. If, for example, a company both publishes the model's picks and also holds inventory of those wines, they could benefit from others buying in (a form of pump-and-dump). It's important that any commercial use has clear policies (like if a platform recommends wines, they disclose if they own those wines or not). As researchers, our role is to present the tool; if we were to spin-off a product, we would adhere to ethical investment advisory standards and possibly regulatory oversight (though wine investing is less regulated than securities).

8.6 Unintended consequences for producers

If investors heavily use such models, producers might start tailoring their winemaking to what the model (and thus the market) prefers, potentially stifling creativity. For instance, if the model rewards wines with certain critic-approved profiles, wineries might feel pressure to conform to those styles, reducing diversity. While speculative, it's worth considering. The flip side: producers who see they score well in such analytics might leverage that in marketing (which is fine if truthful). We should be careful that the model's outputs (like a numeric "investment score") aren't misused to brand wines in a simplistic way (every wine has nuances beyond one score).

In our research dissemination, we will highlight these ethical issues. We'll note that algorithmic predictions should complement, not replace, traditional connoisseurship and due diligence. The idea is to enhance the market's efficiency and access to information, not to concentrate power or manipulate outcomes.

From a commercial perspective, if the research proves fruitful, it could be applied in various ways: an investment advisory tool, a feature in a wine trading platform, or internal use by funds. We will consider the business model implications (selling the software vs. using it internally to run a fund). Each has ethical angles (selling widely vs exclusivity – broad access would democratize the tool, exclusive use could give one party an unfair edge). As academics, our bias is toward sharing knowledge, but we recognize the value of the model might invite commercialization.

In conclusion, we will ensure our research adheres to ethical standards: using data responsibly, avoiding harm to market fairness, and being transparent about the model's capabilities and limits. By doing so, we can help integrate AI into fine wine investment in a way that benefits investors, producers, and the integrity of this niche market.

9 Conclusion & Future Work

9.1 Conclusion

This proposed research will pioneer a comprehensive, multi-modal deep learning approach to fine wine investment prediction. By integrating subregional and vineyard-specific insights with market and qualitative data, we aim to advance beyond traditional wine valuation models. The research's contributions are expected to be multi-fold. Academically, it will bridge the gap between wine economics and machine learning, demonstrating how alternative assets can be modeled with modern AI techniques. Practically, it could yield a tool that helps investors identify promising wines and manage portfolios with greater precision. Our model will encapsulate factors ranging from the macro (market trends, climate shifts) to the micro (terroir, bottle condition), thus offering a holistic assessment of a wine's investment potential. If successful, this work will show that incorporating granular domain knowledge (like vineyard data) into an AI model significantly improves predictions – a lesson that could extend to other collectibles and commodities.

We expect to find that fine-grained features and multi-modal data do improve prediction accuracy and that the model can uncover non-linear interactions (for example, a high critic score *and* limited production together might be required for huge price jumps, which a linear model might miss). We also anticipate providing evidence that fine wine's performance can be forecasted to a useful extent, reinforcing the asset class's viability in portfolios with a scientific backing.

9.2 Future Work

This project opens several avenues for future exploration:

• **Dynamic and live forecasting:** Building on our model, a natural next step is to create a real-time wine price forecasting system. This could involve streaming new auction results or market data into an updated model (perhaps using online learning). One could envision a "wine investment dashboard" that updates projections continuously. Future research might implement reinforcement learning where the model not only predicts but also suggests trading actions (buy/sell/hold) and learns from the results, akin to an algorithmic trading bot for wines.

• **Vineyard mapping and geospatial analysis:** We touched on using vineyard data; future work could expand this by deeply integrating GIS (Geographic Information System) data. For instance, mapping all Burgundy climats with their attributes, and using spatial models to predict how a particular plot's value might change in climate change scenarios. One could combine our price prediction with land value prediction to advise where new vineyards might be smart investments (a sort of "terroir arbitrage"). Additionally, the model could be extended to predict not just price but also optimal drinking windows, which might indirectly affect investment (if a wine is near peak drinkability, some collectors might sell, affecting price).

• **Extension to other alternative assets:** The framework developed (multi-modal deep learning for asset prediction) could be adapted to other collectible markets – e.g., applying a similar model to rare whisky, art, or classic cars. Each of those has analogs (images of the item, descriptions, historical sale data, etc.). A comparative study could be done to see which asset's values are most predictable with AI, or to build a unified model that rates "passion investments" in general.

• **Incorporating consumer sentiment and social media:** Future studies could incorporate data from forums, social media, or search trends to gauge real-time consumer interest in certain wines. Much like how Google Trends are used in stock market prediction, one could use search frequency of "Château X" or sentiment on wine message boards as inputs. Our current model mostly uses expert input; adding consumer sentiment could improve short-term price movement predictions (hype cycles, etc.).

• **Economic and regulatory implications:** Further research could delve into how the emergence of AI-driven investment in wine might influence regulation of the wine trade or tax policy (wine is often treated as a wasting asset for tax in the UK, etc.). As wine investment becomes more data-driven, perhaps new financial instruments (like wine futures or securitized wine funds) will arise. Our model could be used to price derivatives on wine (like options on a wine's future price), which could be a fascinating finance extension.

• **Aging and quality prediction:** Another future angle is to integrate chemical or sensory data of wines. For instance, can we predict how a wine will score in the future (as it ages) using machine learning on its initial release data (like chemical composition)? That goes somewhat beyond price into quality forecasting, but since quality drives price, the two are linked. Partnering with enology labs to get spectroscopy data and then predicting future critic scores would complement our work.

Finally, we note that as with any AI model, continual learning is important. The model can be periodically retrained as new data comes (new vintages, new market shifts). Future work could explore active learning: identifying which new data points (e.g., a new entrant wine) the model is uncertain about, and focusing analysis there.

In conclusion, the proposed research is an ambitious blending of domain expertise and advanced analytics. It promises to yield insights beneficial to investors, researchers, and the wine industry at large. By demonstrating how subregional specifics and multi-modal data improve fine wine investment prediction, we aim to set a precedent for the analysis of other niche asset classes. The knowledge gained will not only improve financial outcomes but also deepen understanding of the fine wine market's mechanics. We look forward to executing this project and contributing both a novel predictive model and a richer scholarly understanding of fine wine as an investment.